

# Mitigating Unfairness in Deep Learning

Guest Lecture – CS 594: Responsible Data Science and Algorithmic Fairness

Vishnu Dasu, Ph.D. student in Computer Sciences

# About Me

- Second year PhD student at Penn State
  - Advised by Prof. Gary Tan
- 2 years industry experience as a security researcher
- Research Interests: Trustworthy AI, Security & Privacy, Applied Cryptography
- Website: <https://vdasu.github.io>
- Contact: vdasu@psu.edu



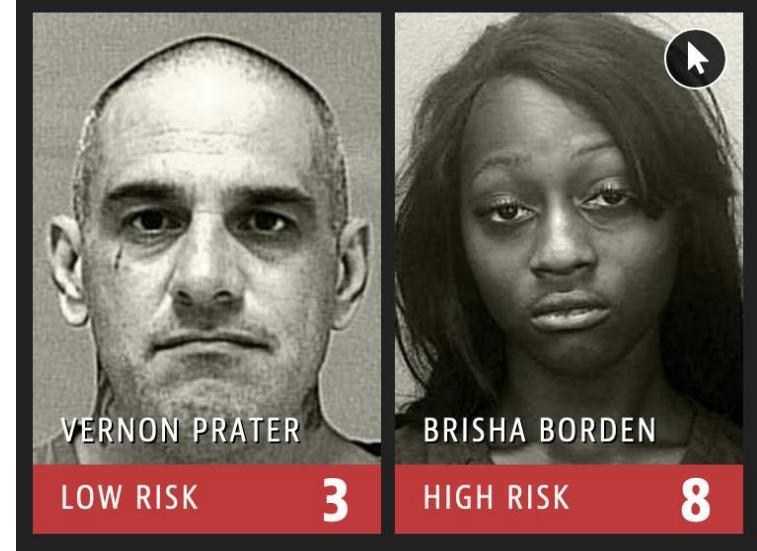
# Fairness in ML Applications



Tax Auditing



Amazon Hiring<sup>1</sup>

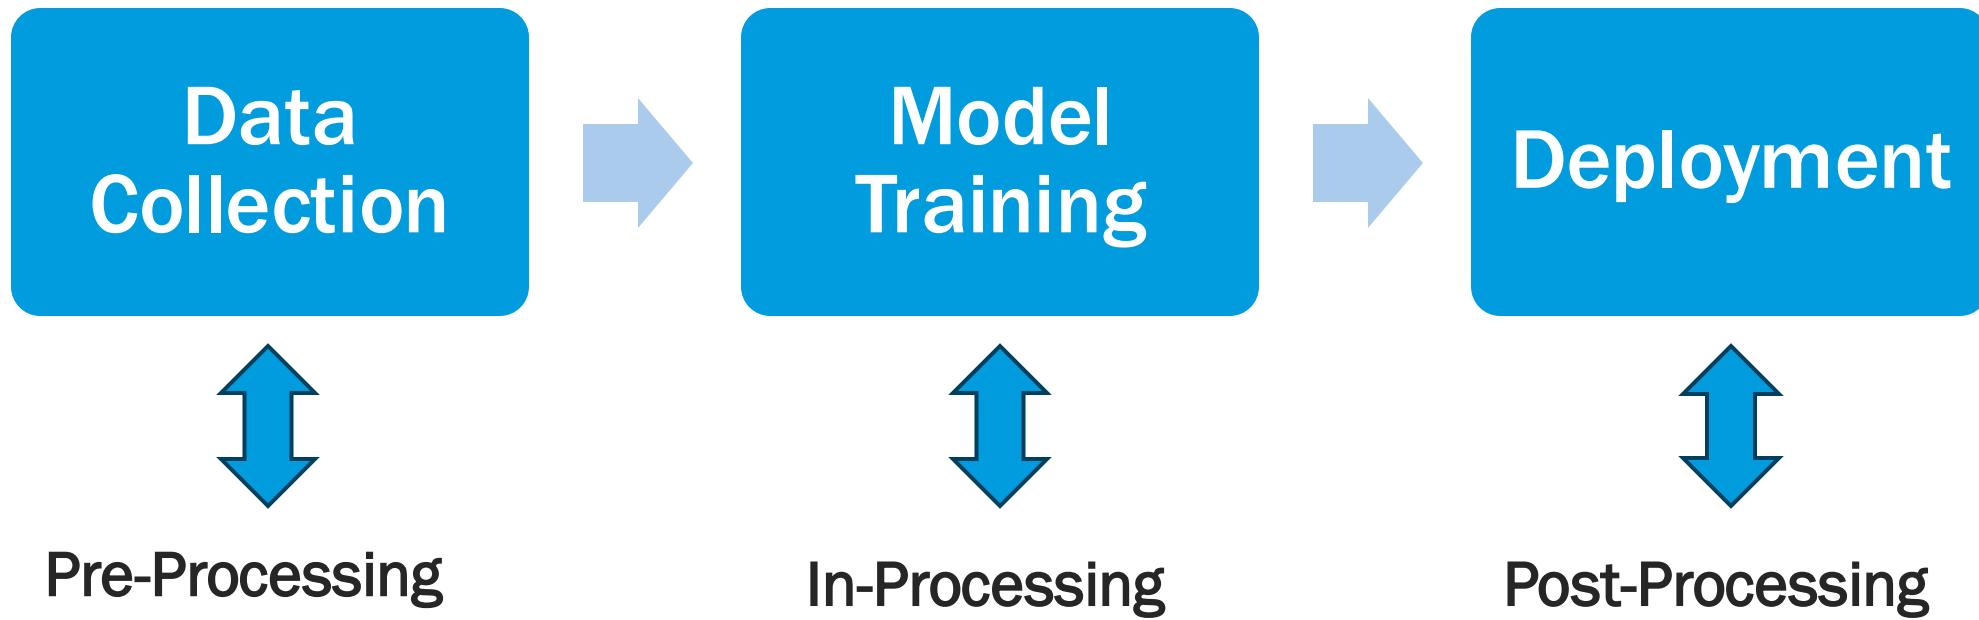


COMPAS<sup>2</sup>

# ML Application Workflow



# Repairing Unfairness in ML



# Repairing Unfairness in ML

- Data collection and training is an expensive and difficult process
- Modifying them is often infeasible
  - Terabytes of data are scraped to train LLMs
  - Cleaning and analyzing the data is impossible
- Can we fix fairness issues after model training is complete?

# Repairing Unfairness in ML

- Data collection and training is an expensive and difficult process
- Modifying them is often infeasible
  - Terabytes of data are scraped to train LLMs
  - Cleaning and analyzing the data is impossible
- Can we fix fairness issues after model training is complete?

Yes!

# Overview

- Fairness in Deep Neural Networks (DNNs)
  - NeuFair: Neural Network Fairness Repair with Dropout [[ISSTA '24](#)]
- Fairness in Large Language Models (LLMs)
  - Attention Pruning: Automated Fairness Repair of Language Models via Surrogate Simulated Annealing [[ICSE '26](#)]

# **NeuFair: Neural Network Fairness Repair with Dropout**

Vishnu Asutosh Dasu, Ashish Kumar, Saeid Tizpaz-Niari, Gang Tan

ISSTA '24

# Problem Statement

Can we repair unfairness in a trained DNN without modifying dataset or re-training?

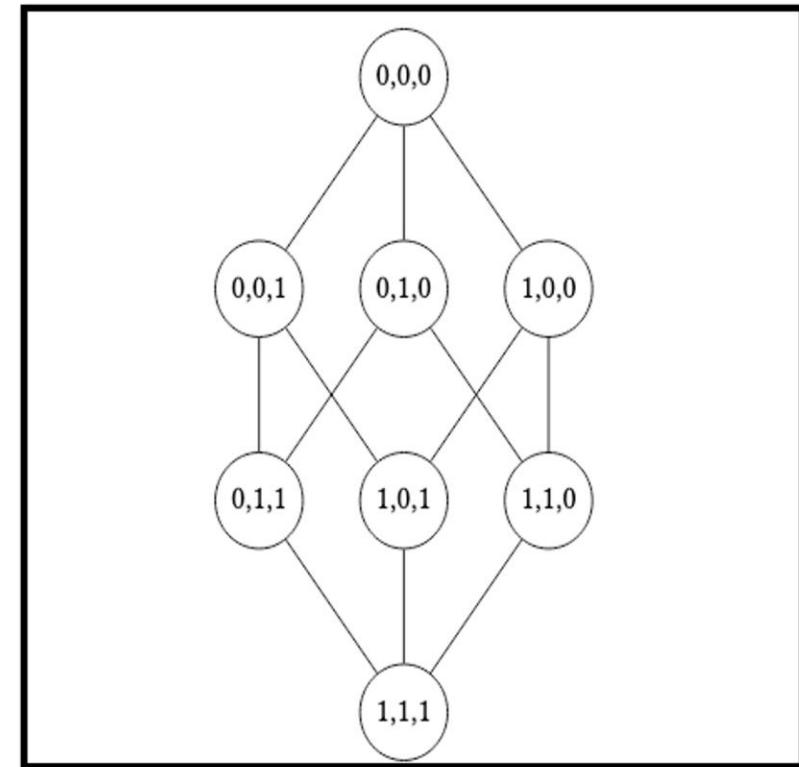
# Key Observation and Idea

**Observation:** Subset of neurons disparately affects fairness

**Idea:** Dropping these neurons after training can improve fairness with negligible loss in utility

# Challenges

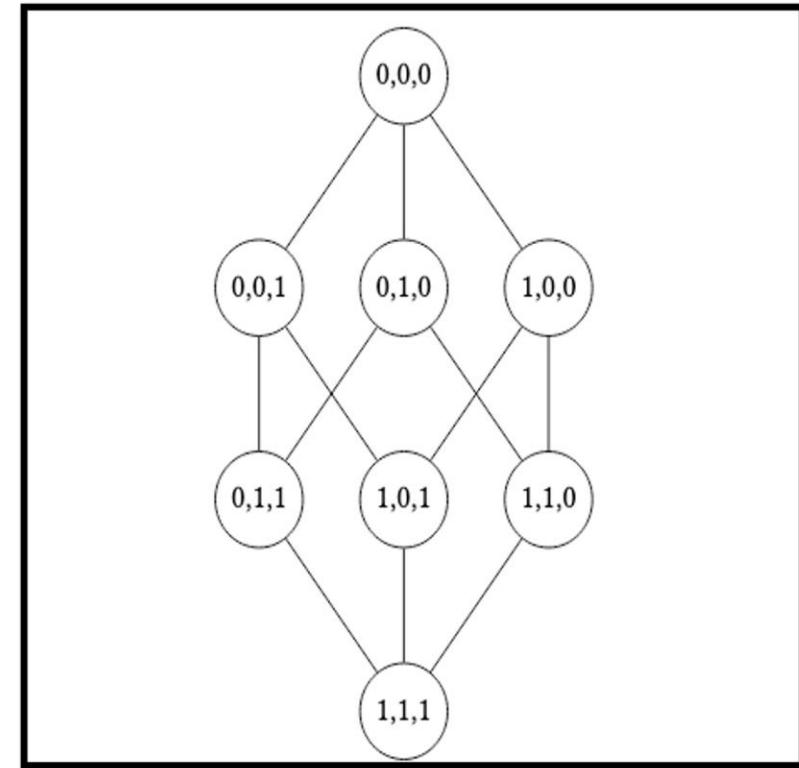
- Exponentially large search space
- DNN with  $N$  neurons has  $2^N$  possible subsets



Search space with  $N = 3$

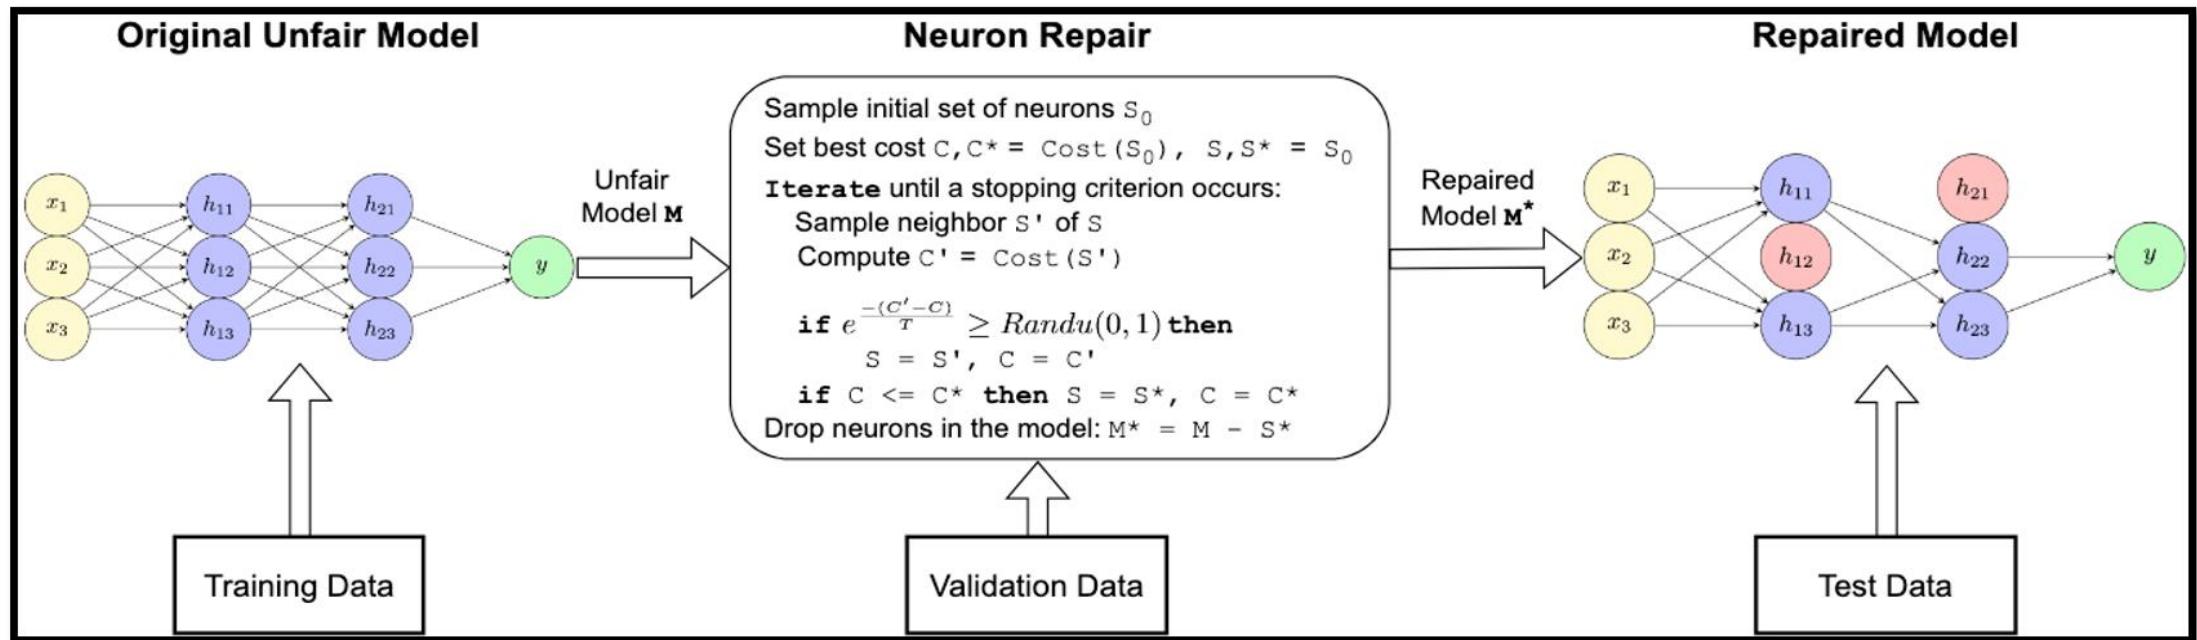
# Solution

- Use randomized algorithms to explore search space
- Two strategies:
  - Simulated Annealing (SA)
  - Random Walk (RW)



Search space with  $N = 3$

# Overview of NeuFair



# Fairness and Utility Definitions

- Fairness:
  - Equalized Odds Difference (EOD) is maximum of difference between true and false positive rates across protected groups

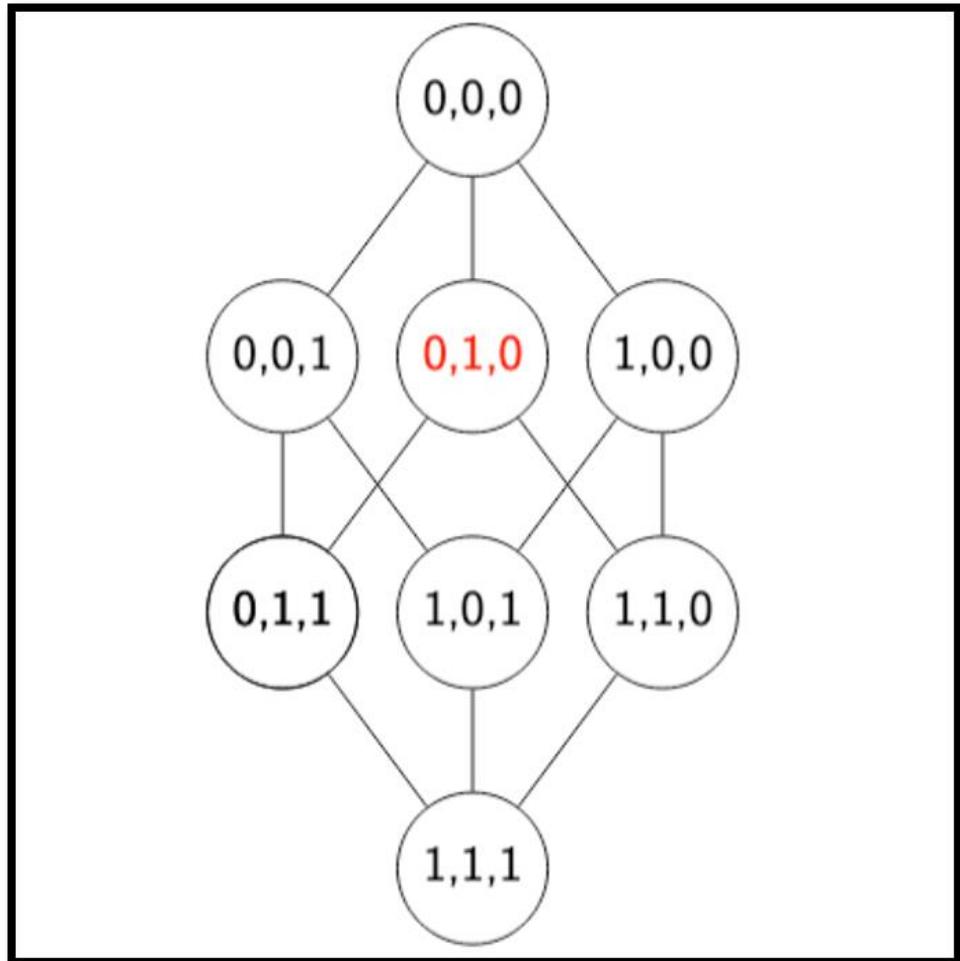
$$EOD := \max(|TPR_A - TPR_B|, |FPR_A - FPR_B|)$$

- Model Utility:
  - F1 Score and Accuracy

$$Acc = \frac{TP + TN}{TP + TN + FP + FN}$$
$$F1 = \frac{2 * TP}{2 * TP + FP + FN}$$

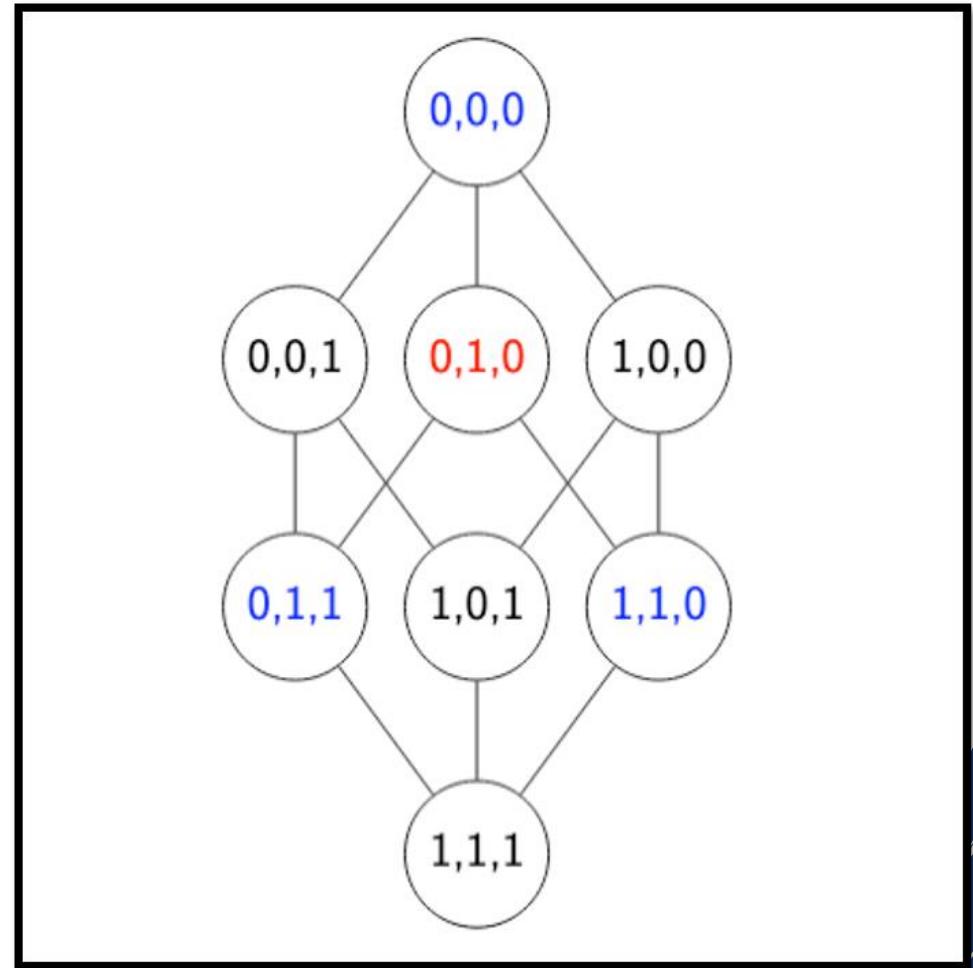
# Methodology

1. Compute cost of current state



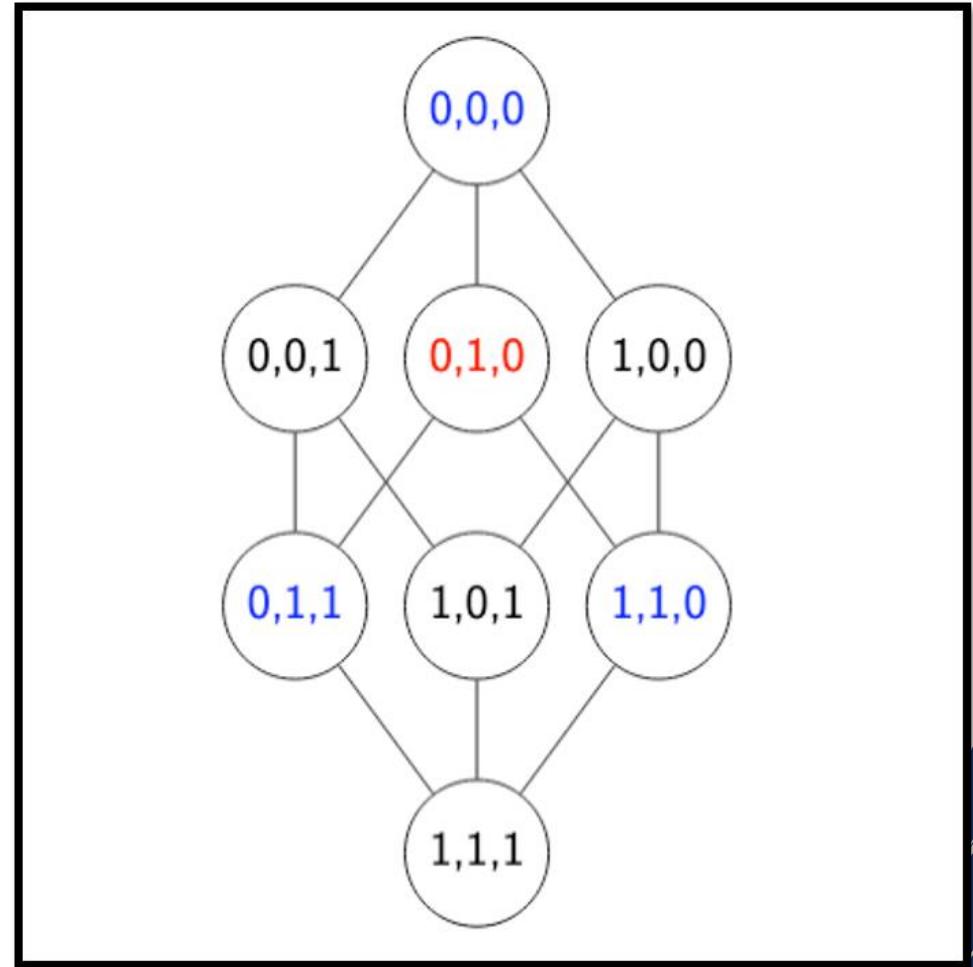
# Methodology

1. Compute cost of current state
2. Sample neighbor and compute cost



# Methodology

1. Compute cost of current state
2. Sample neighbor and compute cost
3. If new cost is better, accept state.  
Else, accept with some probability



# Cost Function

- Minimize EOD and maintain F1 score

1. Compute cost of current state
2. Sample a neighbor state and compute cost
3. If new cost is better, accept state. Else, accept with some probability

# Cost Function

- Minimize EOD and maintain F1 score
- Linear combination of EOD and penalty for poor F1

$$\text{cost}(s) := EOD_s + p \cdot EOD_{s_0} \cdot \mathbb{1}(F1_s < tF1_{s_0})$$

- Fair states are encouraged
- States with low F1 are penalized

1. Compute cost of current state
2. Sample a neighbor state and compute cost
3. If new cost is better, accept state. Else, accept with some probability

# Neighbors of a State

- State stores a subset of neurons that need to be dropped
- Neighbors of a state are all states that add or remove a neuron to its subset
- Hamming Difference is 1
- Pick random neighbor

1. Compute cost of current state
2. Sample a neighbor state and compute cost
3. If new cost is better, accept state. Else, accept with some probability

# Acceptance Criteria

- Always accept good states
- For bad states, two strategies:
  - Simulated Annealing (SA): Accept with probability  $p = e^{-(\Delta \text{cost}/\text{temperature})}$
  - Random Walk (RW): Always accept state i.e. probability  $p = 1$
- SA balances exploration/exploitation
- RW always explores

1. Compute cost of current state
2. Sample a neighbor state and compute cost
3. If new cost is better, accept state. Else, accept with some probability

# NeuFair

Repeat  
Until  
Time Limit

## Algorithm 2: NeuFair to mitigate unfairness in trained neural networks

**Input:** Unfair neural network  $\mathcal{M}$ , Penalty multiplier  $p$ , Threshold multiplier  $t$ , Minimum and maximum number of neurons to drop  $[n_l, n_u]$ , Algorithm Type  $alg\_type$ , Time Limit  $time\_limit$

**Output:** Repaired neural network  $\mathcal{M}_\star$ , Desirable state  $s_\star$ , Best cost  $cost_\star$

```
1  $s \leftarrow random\_state(\mathcal{M}, n_l, n_u)$ 
2  $s_\star, start\_time \leftarrow \phi, curr\_time()$ 
3  $cost \leftarrow compute\_cost(\mathcal{M}, s, p, t)$ 
4  $cost_\star \leftarrow compute\_cost(\mathcal{M}, s_\star, p, t)$ 
5  $T_0 \leftarrow estimate\_temperature(\mathcal{M}, s)$ 
6 while  $curr\_time() - start\_time \leq time\_limit$  do
7    $T \leftarrow update\_temperature(T_0, curr\_time())$ 
8    $s_i \leftarrow generate\_state(s, n_l, n_u)$ 
9    $cost_i \leftarrow compute\_cost(\mathcal{M}, s_i, p, t)$ 
10   $\Delta E \leftarrow cost_i - cost$ 
11  if  $\Delta E \leq 0$  then
12     $cost \leftarrow cost_i$ 
13     $s \leftarrow s_i$ 
14  else if ( $alg\_type == RW$ )  $\vee$  ( $alg\_type == SA$   $\wedge$   $e^{-\Delta E/T} \geq Uniform(0,1)$ ) then
15     $cost \leftarrow cost_i$ 
16     $s \leftarrow s_i$ 
17  if  $cost \leq cost_\star$  then
18     $cost_\star \leftarrow cost_i$ 
19     $s_\star \leftarrow s_i$ 
20  $\mathcal{M}_\star \leftarrow \mathcal{M} \setminus s_\star$ 
21 return  $\mathcal{M}_\star, s_\star, cost_\star$ 
```



Initialization



Compute cost of  
current and new  
state



Transition to  
new state

# Experiments and Results

- RQ1: How effective are randomized algorithms in repairing unfairness?
- RQ2: Can dropout improve fairness and utility together?
- RQ3: What are the design considerations of search algorithms?
- RQ4: How does NeuFair perform against SOTA post-processing algorithms?

# Experiments and Results

- Seven settings with five different datasets
- Repeat with 10 random seeds each, 1 hr timeout
- Penalty multiplier = 3.0
- F1 threshold = 0.98

$$C(s) = EOD_s + 3.0 \cdot EOD_{s_0} \cdot \mathbb{1}(F1_s < (0.98 \cdot F1_{s_0}))$$

- Restricted search space:
  - Min. 2 neurons
  - Max. 20-40% of DNN
- 60%-20%-20% Train-Validation-Test split

# RQ1: Effective of randomized algorithms

Datasets	Original Model			Simulated Annealing		
	EOD	F1	Accuracy	EOD	F1	Accuracy
Adult (Sex)	11.639% $\pm$ 2.326	0.667 $\pm$ 0.008	0.851 $\pm$ 0.003	7.259% $\pm$ 1.697	0.652 $\pm$ 0.01	0.849 $\pm$ 0.004
Adult (Race)	8.251% $\pm$ 3.195			4.976% $\pm$ 1.816	0.656 $\pm$ 0.008	0.849 $\pm$ 0.004
COMPAS (Sex)	2.522% $\pm$ 0.817	0.967 $\pm$ 0.004	0.969 $\pm$ 0.004	2.921% $\pm$ 1.446	0.954 $\pm$ 0.08	0.957 $\pm$ 0.008
COMPAS (Race)	2.96% $\pm$ 1.088			2.239% $\pm$ 1.003	0.954 $\pm$ 0.005	0.957 $\pm$ 0.004
Bank	14.665% $\pm$ 2.114	0.553 $\pm$ 0.004	0.84 $\pm$ 0.003	7.257% $\pm$ 3.533	0.537 $\pm$ 0.014	0.888 $\pm$ 0.01
Default	8.962% $\pm$ 1.772	0.53 $\pm$ 0.006	0.769 $\pm$ 0.007	2.749% $\pm$ 0.827	0.519 $\pm$ 0.006	0.79 $\pm$ 0.015
MEPS16	20.641% $\pm$ 2.527	0.533 $\pm$ 0.01	0.788 $\pm$ 0.009	8.426% $\pm$ 2.311	0.507 $\pm$ 0.02	0.853 $\pm$ 0.005

Simulated Annealing on Test split

# RQ1: Effective of randomized algorithms

Datasets	Original Model			Random Walk		
	EOD	F1	Accuracy	EOD	F1	Accuracy
Adult (Sex)	11.639% $\pm$ 2.326	0.667 $\pm$ 0.008	0.851 $\pm$ 0.003	7.358% $\pm$ 1.063	0.652 $\pm$ 0.01	0.849 $\pm$ 0.004
Adult (Race)	8.251% $\pm$ 3.195			4.785% $\pm$ 2.085	0.658 $\pm$ 0.009	0.849 $\pm$ 0.003
COMPAS (Sex)	2.522% $\pm$ 0.817	0.967 $\pm$ 0.004	0.969 $\pm$ 0.004	2.233% $\pm$ 1.022	0.954 $\pm$ 0.007	0.957 $\pm$ 0.006
COMPAS (Race)	2.96% $\pm$ 1.088			2.159% $\pm$ 1.07	0.955 $\pm$ 0.004	0.958 $\pm$ 0.004
Bank	14.665% $\pm$ 2.114	0.553 $\pm$ 0.004	0.84 $\pm$ 0.003	7.595% $\pm$ 2.733	0.548 $\pm$ 0.008	0.881 $\pm$ 0.008
Default	8.962% $\pm$ 1.772	0.53 $\pm$ 0.006	0.769 $\pm$ 0.007	3.124% $\pm$ 0.937	0.523 $\pm$ 0.005	0.79 $\pm$ 0.015
MEPS16	20.641% $\pm$ 2.527	0.533 $\pm$ 0.01	0.788 $\pm$ 0.009	9.86% $\pm$ 2.623	0.513 $\pm$ 0.018	0.851 $\pm$ 0.007

Random Walk on Test split

## RQ2: Improve Fairness and Utility together

Datasets	Original Model			Simulated Annealing		
	EOD	F1	Accuracy	EOD	F1	Accuracy
Adult (Sex)	$11.639\% \pm 2.326$	$0.667 \pm 0.008$	$0.851 \pm 0.003$	$7.259\% \pm 1.697$	$0.652 \pm 0.01$	$0.849 \pm 0.004$
Adult (Race)	$8.251\% \pm 3.195$			$4.976\% \pm 1.816$	$0.656 \pm 0.008$	$0.849 \pm 0.004$
COMPAS (Sex)	$2.522\% \pm 0.817$	$0.967 \pm 0.004$	$0.969 \pm 0.004$	$2.921\% \pm 1.446$	$0.954 \pm 0.08$	$0.957 \pm 0.008$
COMPAS (Race)	$2.96\% \pm 1.088$			$2.239\% \pm 1.003$	$0.954 \pm 0.005$	$0.957 \pm 0.004$
Bank	$14.665\% \pm 2.114$	$0.553 \pm 0.004$	$0.84 \pm 0.003$	$7.257\% \pm 3.533$	$0.537 \pm 0.014$	$0.888 \pm 0.01$
Default	$8.962\% \pm 1.772$	$0.53 \pm 0.006$	$0.769 \pm 0.007$	$2.749\% \pm 0.827$	$0.519 \pm 0.006$	$0.79 \pm 0.015$
MEPS16	$20.641\% \pm 2.527$	$0.533 \pm 0.01$	$0.788 \pm 0.009$	$8.426\% \pm 2.311$	$0.507 \pm 0.02$	$0.853 \pm 0.005$

Simulated Annealing on Test split

## RQ2: Improve Fairness and Utility together

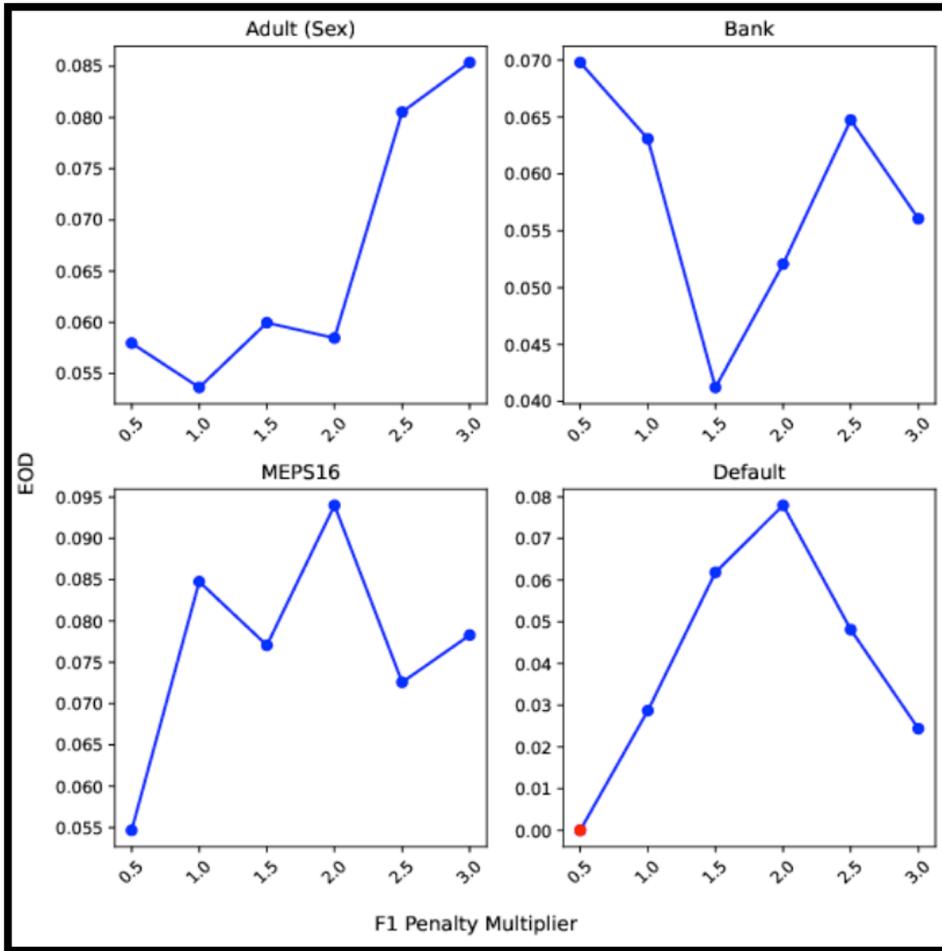
Datasets	Original Model			Random Walk		
	EOD	F1	Accuracy	EOD	F1	Accuracy
Adult (Sex)	11.639% $\pm$ 2.326	0.667 $\pm$ 0.008	0.851 $\pm$ 0.003	7.358% $\pm$ 1.063	0.652 $\pm$ 0.01	0.849 $\pm$ 0.004
Adult (Race)	8.251% $\pm$ 3.195			4.785% $\pm$ 2.085	0.658 $\pm$ 0.009	0.849 $\pm$ 0.003
COMPAS (Sex)	2.522% $\pm$ 0.817	0.967 $\pm$ 0.004	0.969 $\pm$ 0.004	2.233% $\pm$ 1.022	0.954 $\pm$ 0.007	0.957 $\pm$ 0.006
COMPAS (Race)	2.96% $\pm$ 1.088			2.159% $\pm$ 1.07	0.955 $\pm$ 0.004	0.958 $\pm$ 0.004
Bank	14.665% $\pm$ 2.114	0.553 $\pm$ 0.004	0.84 $\pm$ 0.003	7.595% $\pm$ 2.733	0.548 $\pm$ 0.008	0.881 $\pm$ 0.008
Default	8.962% $\pm$ 1.772	0.53 $\pm$ 0.006	0.769 $\pm$ 0.007	3.124% $\pm$ 0.937	0.523 $\pm$ 0.005	0.79 $\pm$ 0.015
MEPS16	20.641% $\pm$ 2.527	0.533 $\pm$ 0.01	0.788 $\pm$ 0.009	9.86% $\pm$ 2.623	0.513 $\pm$ 0.018	0.851 $\pm$ 0.007

Random Walk on Test split

## RQ3: Design Considerations of Search Algorithms

- Increasing F1 threshold multiplier may decrease fairness
- Increasing search space (min/max no: of neurons) and may increase fairness
- Increasing time out may increase fairness
- F1 penalty multiplier has non-trivial effect

## RQ3: Design Considerations of Search Algorithms



Effect of Penalty Multiplier on EOD

$$\text{cost}(s) := EOD_s + p \cdot EOD_{s_0} \cdot \mathbb{I}(F1_s < tF1_{s_0})$$

SA acceptance probability  $p = e^{-(\Delta \text{cost}/\text{temperature})}$

## RQ4: Comparison to SOTA post-processing algorithms

Datasets	EOD		
	Original	DICE	NeuFair (SA)
Adult (Sex)	11.639% $\pm$ 2.326	10.453% $\pm$ 2.266	7.259% $\pm$ 1.697
Adult (Race)	8.251% $\pm$ 2.236	8.092% $\pm$ 3.253	4.976% $\pm$ 1.816
COMPAS (Sex)	2.522% $\pm$ 0.817	3.229% $\pm$ 0.774	2.921% $\pm$ 1.446
COMPAS (Race)	2.96% $\pm$ 1.088	2.964% $\pm$ 1.088	2.239% $\pm$ 1.003
Bank	14.665% $\pm$ 2.114	12.205% $\pm$ 2.731	7.257% $\pm$ 3.533
Default	8.962% $\pm$ 1.772	5.845% $\pm$ 1.816	2.749% $\pm$ 0.827
MEPS16	20.641% $\pm$ 2.527	19.204% $\pm$ 2.592	8.426% $\pm$ 2.311

NeuFair compared to DICE (ICSE 23')

# Conclusion

- A subset of neurons disparately affects fairness
- Deterministic dropout after training can improve fairness without loss in utility
- Randomized algorithms are effective for neuron dropout

# Attention Pruning: Automated Fairness Repair of Language Models via Surrogate Simulated Annealing

Vishnu Asutosh Dasu, Md Rafi Rashid, Vipul Gupta, Saeid Tizpaz-Niari, Gang Tan

ICSE '26

# Problem Statement

Can we extend the idea of *NeuFair* to repair fairness in pre-trained LLMs?

# Challenges

1. **Size**: LLMs contain billions of parameters
  - NeuFair experiments with small DNNs with 100s of neurons
2. **Time**: LLMs have high-inference time
  - Exploring each state takes several minutes
  - SA needs to explore 1000s of states in the exponential search space

# Fairness Definition

- Bias *HolisticBias* dataset:
  - A collection of prompts belonging to different groups  $G$  (e.g. race, gender, sex, etc.)
  - Each group  $G$  contains sub-groups  $g$  (gender: man, woman, transgender, etc.)
  - Bias is measured as differences in model behavior (e.g., toxicity) across sub-groups within the same group

$$\text{bias}_{\Theta}(G) = \sum_{g \in G} |T_G - T_g|$$

where  $T_g = \frac{1}{|D_g|} \sum_{x \in D_g} \text{tox}_{\Theta}(x)$

and  $T_G = \frac{1}{|G|} \sum_{g \in G} T_g$

# Utility Definition

- *Perplexity (PPL)* measures how well a model  $\Theta$  predicts a sequence  $y$
- It is defined as the exponential of the average negative log-likelihood the model assigns to each token
- We use the WikiText-2 dataset to measure perplexity

$$PPL_{\Theta}(y) = \exp \left( -\frac{1}{n} \sum_{i=1}^n \log(\Theta(y_i | y_1, \dots, y_{i-1})) \right)$$

# Dealing with the Size of LLMs

# Attention in LLMs

- Scaled dot-product attention (SDPA) operates on Query, Key, and Value vectors

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

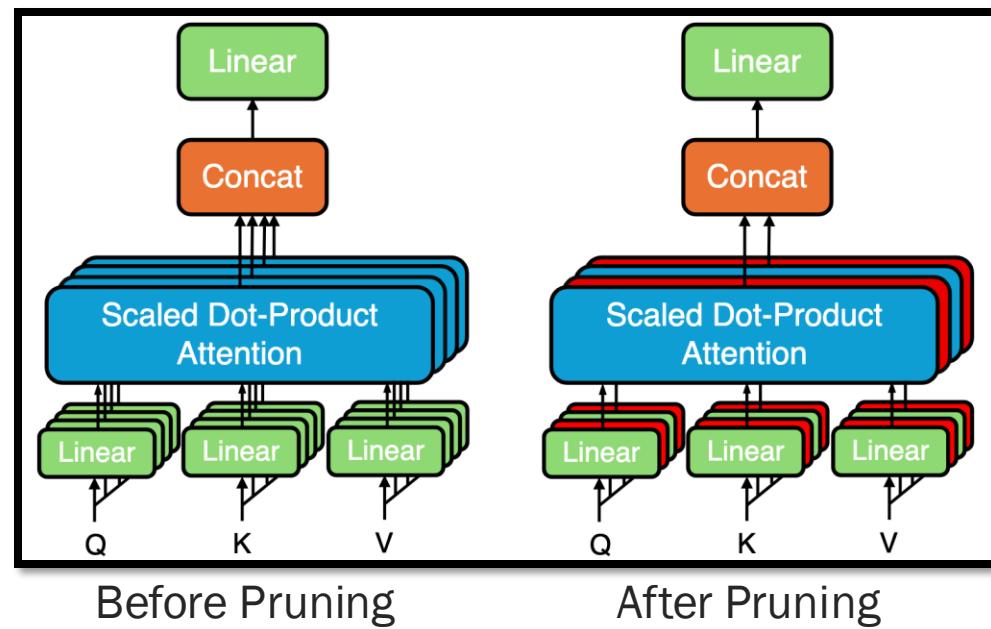
- Transformer-based LLMs consist of stacked blocks of Multi Head Attention (MHA) in each block

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_n)W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

# Pruning Attention Heads

- Several works have shown that LLMs are over-parameterized
- Pruning a subset of head of attention heads can improve efficiency with minimal loss in performance



# Pruning Attention Heads

- Can we prune attention heads to improve fairness?
  - Less granular than infeasible neuron-level pruning and faster

# Pruning Attention Heads

- Can we prune attention heads to improve fairness?
  - Less granular than infeasible neuron-level pruning and faster
- *FASP Algorithm*: Fairness-Aware Structured Pruning in Transformers [Zayed et. al, AAAI 2024]<sup>1</sup>

# The FASP algorithm

1. Prune each attention head one at a time while keeping others intact

# The FASP algorithm

1. Prune each attention head one at a time while keeping others intact
2. Calculate the change in fairness and utility after pruning

# The FASP algorithm

1. Prune each attention head one at a time while keeping others intact
2. Calculate the change in fairness and utility after pruning
3. Rank attention heads by magnitude of change

# The FASP algorithm

1. Prune each attention head one at a time while keeping others intact
2. Calculate the change in fairness and utility after pruning
3. Rank attention heads by magnitude of change
4. Define *critical* heads: Heads that are important for utility

# The FASP algorithm

1. Prune each attention head one at a time while keeping others intact
2. Calculate the change in fairness and utility after pruning
3. Rank attention heads by magnitude of change
4. Define *critical* heads: Heads that are important for utility
5. Prune a subset of *non-critical* heads that improve fairness the most

# The FASP algorithm: Issues

- Effect of pruning attention heads is non-linear!
  - Heads interact through residual connections and layer norms
- E.g. If pruning the 1<sup>st</sup> and 3<sup>rd</sup> individually improves fairness, pruning them together may not

# The FASP algorithm: Issues

- Effect of pruning attention heads is non-linear!
  - Heads interact through residual connections and layer norms
- E.g. If pruning the 1<sup>st</sup> and 3<sup>rd</sup> individually improves fairness, pruning them together may not
- **Solution:** Use Simulated Annealing to consider *all* possible subsets of attention heads!

# Dealing with the Inference Time of LLMs

# Efficiency of Simulated Annealing (SA) with LLMs

- SA for pruning attention heads solves the non-linearity problem
- However, SA needs to explore thousands of subsets of attention heads and inferring the fairness/utility of each state takes several minutes
- E.g. One round of inference on LLama-2-7B takes 13 minutes on RTX A6000
- How do we scale?

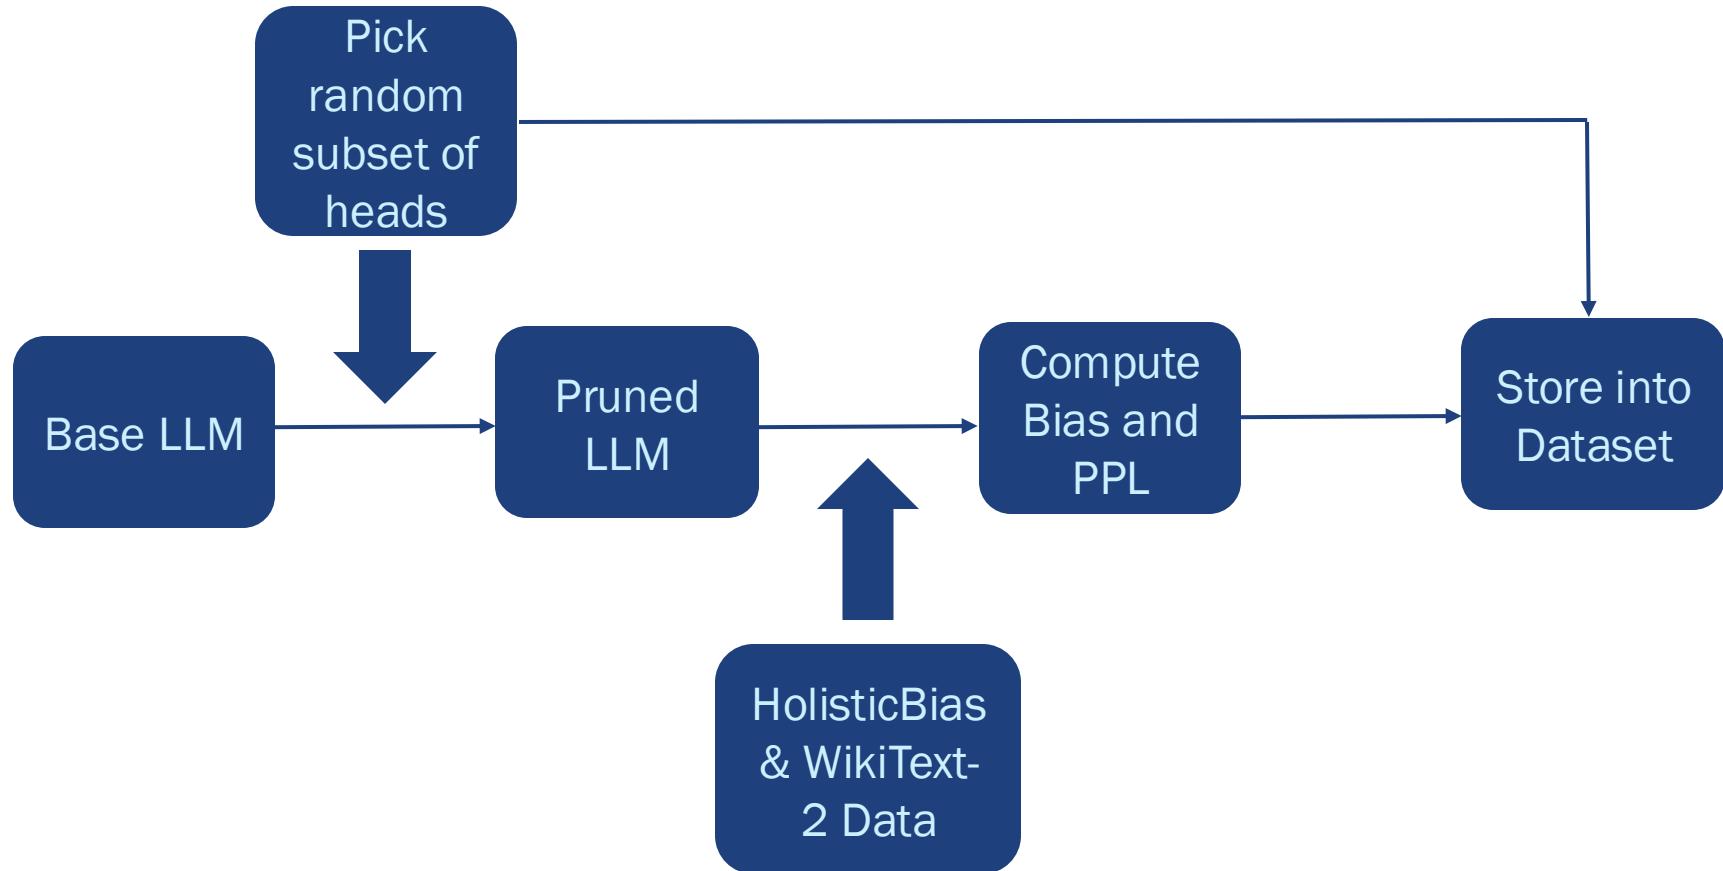
# Approximating Fairness/Utility

- Instead of computing the real fairness/utility of every state, can we approximate it?

# Approximating Fairness/Utility

- Instead of computing the real fairness/utility of every state, can we approximate it?
- **Solution:** Train DNNs to predict what the fairness/utility after pruning a subset of attention heads would be

# Surrogate DNN Training: Dataset Creation



# Examples from Dataset

<b>Model</b>	<b>Attention Head Configuration</b>	<b>Bias</b>	<b>Perplexity</b>
Llama-2-7B	0, 1, 0, 0, 0, ..., 0, 0, 0, 0, 0	0.295	20.75
	1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0	0.323	8.0
GPT-J-6B	1, 1, 0, 0, 0, ..., 1, 0, 1, 0, 1	0.487	15.461
	0, 0, 1, 0, 0, ..., 0, 1, 0, 0, 1	0.428	13.383

# Surrogate DNN training

- Train two DNNs,  $\theta_{bias}$  and  $\theta_{ppl}$ , to predict the bias and perplexity
- The input to the DNN is a bit-vector representing a subset of attention heads
- DNNs are trained for regression using MSE loss

---

**Algorithm 1:** Training surrogate DNNs to capture the effect of pruning attention heads on the bias and perplexity.

---

**Input:** Language Model  $\Theta$ , Prompts for bias  $\mathcal{D}_{bias}$ , Text data for perplexity  $\mathcal{D}_{ppl}$ , Min. and max. number of attention heads to drop  $[n_l, n_u]$ , Fraction of dataset to use for bias and perplexity  $[\eta_{bias}, \eta_{ppl}]$ , and Time Limit *time\_limit*.

**Output:** Trained DNNs  $\theta_{bias}$  and  $\theta_{ppl}$  that predict the bias and perplexity of  $\Theta$  after dropping attention heads.

```
1  $\mathcal{B}, \mathcal{P} \leftarrow \emptyset, \emptyset$ 
2 while  $curr\_time() - start\_time \leq time\_limit$  do
3    $s \leftarrow random\_heads(\Theta, n_l, n_u)$ 
4    $\Theta' \leftarrow prune\_heads(\Theta, s)$ 
5    $S_{bias} \subseteq \mathcal{D}_{bias}$  such that  $|S_{bias}| = \eta_{bias} |\mathcal{D}_{bias}|$ 
6    $S_{ppl} \subseteq \mathcal{D}_{ppl}$  such that  $|S_{ppl}| = \eta_{ppl} |\mathcal{D}_{ppl}|$ 
7    $bias \leftarrow compute\_bias(\Theta'(S_{bias}))$ 
8    $ppl \leftarrow compute\_ppl(\Theta'(S_{ppl}))$ 
9    $\mathcal{B} = \mathcal{B} \cup \{(s, bias)\}$ 
10   $\mathcal{P} = \mathcal{P} \cup \{(s, ppl)\}$ 
11  $\theta_{bias} = \arg \min_{\theta} \frac{1}{|\mathcal{B}|} \sum_{(s,y) \in \mathcal{B}} \|\theta(s) - y\|^2$ 
12  $\theta_{ppl} = \arg \min_{\theta} \frac{1}{|\mathcal{P}|} \sum_{(s,y) \in \mathcal{P}} \|\theta(s) - y\|^2$ 
```

---

# Surrogate Simulated Annealing

- Problem is now reduced to finding best input to both DNNs
- **Significantly** scales up SA with LLMs
  - E.g. 2,260,000x speed-up for LLama-2-7b



# Cost Function

- Cost function of SA is weighted combination of bias and PPL DNN
- $\epsilon$  controls the bias vs. perplexity trade off
  - Higher epsilon trades off better bias for worse perplexity

$$\text{cost}(s) := \epsilon \cdot \theta_{bias}(s) + (1 - \epsilon) \cdot \theta_{PPL}(s)$$

# Attention Pruning Algorithm

**Algorithm 2:** AP: Surrogate Simulated Annealing for Fairness-Aware Attention Head Pruning.

**Input:** Unfair LLM  $\Theta$ , DNNs to predict bias and perplexity  $[\theta_{bias}, \theta_{ppl}]$ , min. and max. number of attention heads to drop  $[n_l, n_u]$ , and Timeout  $time\_limit$

**Output:** Repaired LLM  $\Theta_\star$ , Ideal state  $s_\star$ , Best cost  $cost_\star$

```
1  $s \leftarrow random\_state(\Theta, n_l, n_u)$ 
2  $s_\star, start\_time \leftarrow [0, 0, \dots, 0, 0], curr\_time()$ 
3  $cost \leftarrow \theta_{bias}(s) + \theta_{ppl}(s)$ 
4  $cost_\star \leftarrow \theta_{bias}(s_\star) + \theta_{ppl}(s_\star)$ 
5  $T_0 \leftarrow estimate\_temperature(\theta_{bias}, \theta_{ppl}, s)$ 
6 while  $curr\_time() - start\_time \leq time\_limit$  do
7    $T \leftarrow update\_temperature(T_0, curr\_time())$ 
8    $s_i \leftarrow generate\_state(s, n_l, n_u)$ 
9    $cost_i \leftarrow \theta_{bias}(s_i) + \theta_{ppl}(s_i)$ 
10   $\Delta E \leftarrow cost_i - cost$ 
11  if  $\Delta E \leq 0$  then
12     $cost \leftarrow cost_i$ 
13     $s \leftarrow s_i$ 
14  else if  $e^{-\Delta E/T} \geq Uniform(0,1)$  then
15     $cost \leftarrow cost_i$ 
16     $s \leftarrow s_i$ 
17  if  $cost \leq cost_\star$  then
18     $cost_\star \leftarrow cost_i$ 
19     $s_\star \leftarrow s_i$ 
20  $\Theta_\star \leftarrow prune\_heads(\Theta, s_\star)$ 
21 return  $\Theta_\star, s_\star, cost_\star$ 
```

Repeat  
Until  
Time Limit



Initialization after  
training DNNs

Compute cost using  
DNNs

Transition to  
new state

Finally prune  
LLM

# Experiments and Results

- RQ1: How effective are surrogate DNNs at predicting bias/perplexity?
- RQ2: How does Attention Pruning compare to SOTA?
- RQ3: What are the design considerations of Attention Pruning?
- RQ4: Can Attention Pruning generalize beyond Gender bias?

# Experiments and Results

- Six different LLMs evaluated against Gender bias from HolisticBias
- Repeat with 3 different seeds
- Standard pre-processing to scale down bias/PPL in [0,1]
- Time limit of 3hrs

## RQ1: Effectiveness of Surrogate DNNs

Model	$\theta_{bias}$ MSE	$\theta_{ppl}$ MSE
Distilgpt-2	0.0038	0.0005
GPT-2	0.004	0.004
GPT-Neo-125M	0.007	0.007
GPT-Neo-1.3B	0.0049	0.026
GPT-J-6B	0.0073	0.0048
Llama-2-7B	0.0046	0.010

MSE of Trained Surrogate DNNs

## RQ1: Effectiveness of Surrogate DNNs

Model	cost = $\theta_{bias}(s)$		cost = $\theta_{ppl}(s)$		cost = $\epsilon \cdot \theta_{bias}(s) + (1 - \epsilon) \cdot \theta_{ppl}(s)$	
	Bias	PPL	Bias	PPL	Bias	PPL
Distilgpt-2	0.275	211.61	0.415	65.28	0.285	74.981
GPT-2	0.245	80.694	0.415	43.533	0.233	52.071
GPT-Neo 125M	0.196	20226574.0	0.35	39.377	0.236	41.912
GPT-Neo 1.3B	0.249	21.7	0.42	18.323	0.282	18.5
GPT-J 6B	0.276	14.492	0.38	12.258	0.275	13.17
Llama-2 7B	0.37	7.5	0.405	6.688	0.317	7.219

Effect of using only one surrogate DNN with SA

## RQ2: Comparison against state-of-the-art

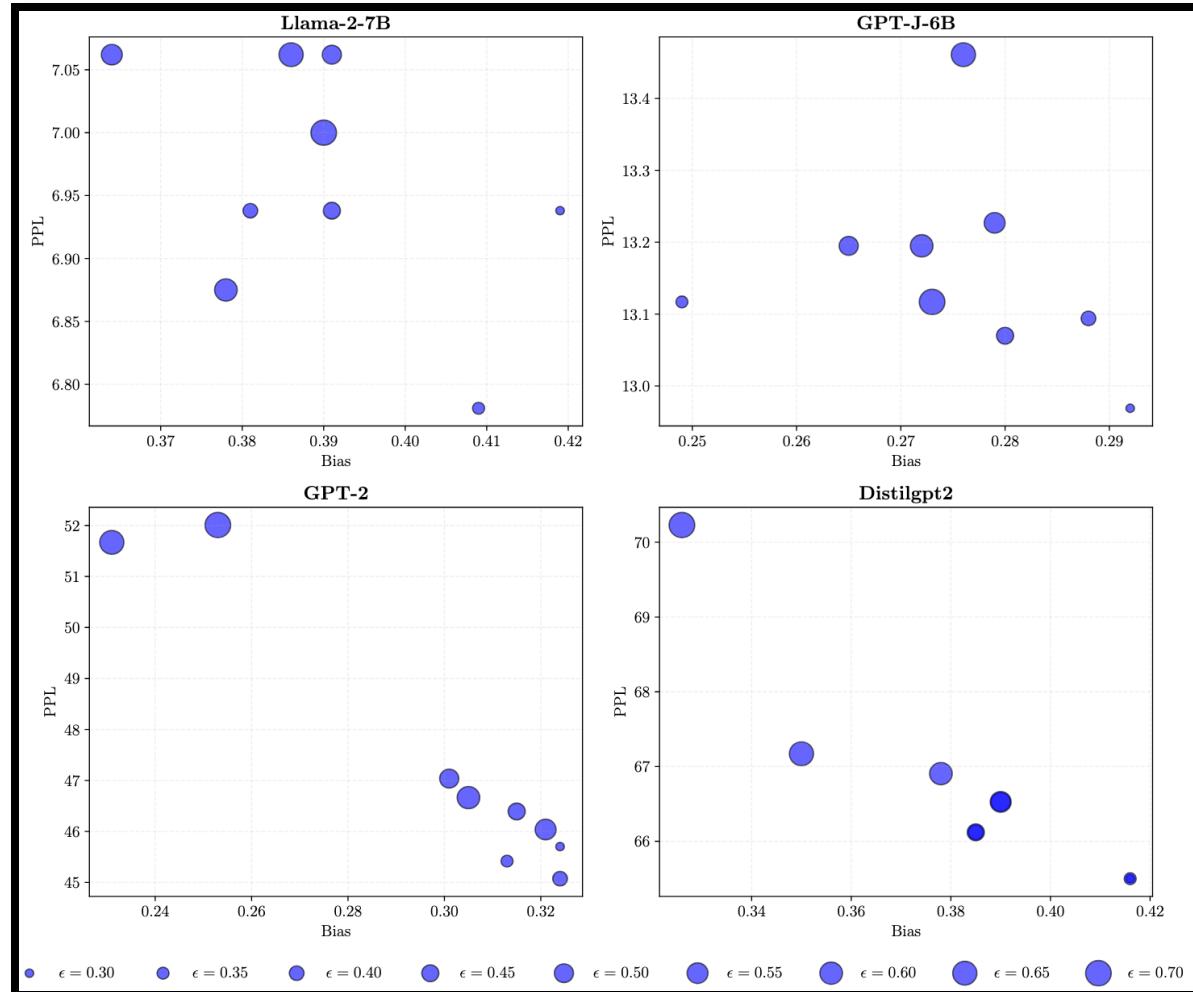
Model	Baseline		AP ( $\epsilon = 0.5$ )		FASP [72] ( $\gamma = 0.3$ )	
	Bias	PPL	Bias	PPL	Bias	PPL
Distilgpt-2	$0.428 \pm 0.028$	65.325	$0.288 \pm 0.002 (\epsilon = 0.88, \eta = 0.2)$	<b>74.891</b>	$0.31 \pm 0.013 (\alpha = 0.2)$	78.212
GPT-2	$0.402 \pm 0.01$	43.588	$0.255 \pm 0.02 (\epsilon = 0.8, \eta = 0.2)$	<b>52.071</b>	$0.27 \pm 0.018 (\alpha = 0.2)$	58.125
GPT-Neo-125M	$0.399 \pm 0.008$	35.628	$0.241 \pm 0.004 (\eta = 0.1)$	<b>41.912</b>	$0.221 \pm 0.004 (\alpha = 0.1)$	47.237
GPT-Neo-1.3B	$0.435 \pm 0.001$	17.422	$0.285 \pm 0.003 (\eta = 0.05)$	<b>18.548</b>	$0.339 \pm 0.012 (\alpha = 0.16, \gamma = 0.6)$	19.391
GPT-J-6B	$0.446 \pm 0.013$	11.695	$0.264 \pm 0.01 (\eta = 0.1)$	<b>13.17</b>	$0.288 \pm 0.005 (\alpha = 0.18)$	13.227
Llama-2-7B	$0.4 \pm 0.006$	6.781	$0.316 \pm 0.003 (\eta = 0.05)$	7.219	$0.342 \pm 0.004 (\alpha = 0.06)$	<b>6.781</b>

Overall, AP finds states with better bias and PPL

## RQ3: Design Considerations of Attention Pruning

- AP has several hyperparameters that need to tuned
- Two important ones:
  - $\epsilon$  in cost function
  - Running Time

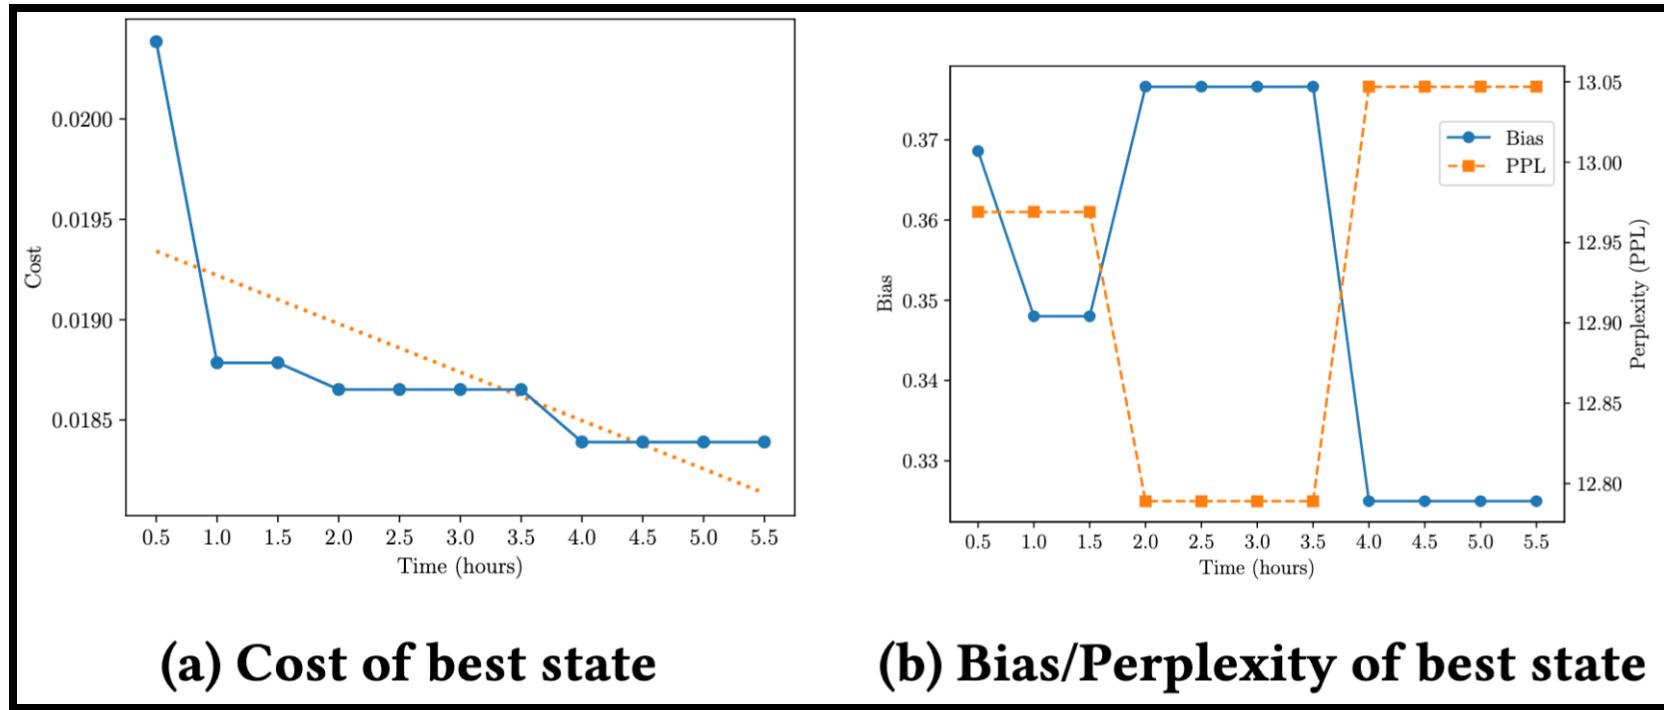
## RQ3: Design Considerations of Attention Pruning



$$\text{cost}(s) := \epsilon \cdot \theta_{bias}(s) + (1 - \epsilon) \cdot \theta_{PPL}(s)$$

Higher values of  $\epsilon$  yield better bias  
at the expense of PPL

## RQ3: Design Considerations of Attention Pruning



Effect of running time on cost and bias/PPL

## RQ4: Effect of reducing gender bias on other biases

Model	Race		Nationality		Sexual Orientation		Age	
	Baseline	AP	Baseline	AP	Baseline	AP	Baseline	AP
Distilgpt-2	0.529	0.395	0.288	0.225	0.737	0.57	0.054	0.036
GPT-2	0.518	0.357	0.301	0.217	0.632	0.512	0.058	0.025
GPT-Neo-125M	0.448	0.33	0.23	0.163	0.792	0.494	0.044	0.027
GPT-Neo-1.3B	0.463	0.374	0.226	0.201	0.672	0.445	0.078	0.048
GPT-J-6B	0.496	0.369	0.258	0.201	0.614	0.471	0.069	0.024
Llama-2-7B	0.458	0.385	0.252	0.217	0.612	0.49	0.057	0.033

Reducing Gender bias also reduces  
other biases

# Conclusion

- Surrogate DNNs are effective at estimating the effect of pruning attention heads on bias/PPL
- Surrogate Simulated Annealing is effective at considering non-linear relationships between attention heads

# Future Works

- Surrogate DNN idea is very interesting and powerful
  - Can we extend to other areas like robustness, privacy, etc.?
- Data collection process to train DNNs is expensive
  - Several hours to collect 1000s of samples
  - Can we find more efficient ways to train surrogate models?
- Using genetic algorithms instead of SA

# Thank You!